Cheng Li

1308 W. Main St., Urbana, Illinois 61801

cli99.netlify.com | cli99@illinois.edu | github.com/rai-project | 4194967797

OBJECTIVE

Full-time position in research and engineering

RESEARCH INTEREST

My research lies in the field of GPU-accelerated applications, with an emphasis on Deep Learning (DL). My work has focused on understanding and optimizing Deep Learning workloads. In the process, I have developed several open-source tools to benchmark, profile, and summarize Deep Learning training and inference across hardware and software stacks.

EDUCATION

University of Illinois Urbana-Champaign — **Champaign, IL**
Ph.D. in Computer Science — *Expected August 2020*
GPA: 3.95/4.0

Thesis: Performance Benchmarking, Analysis and Optimization of Deep Learning Inference

University of Michigan — **Ann Arbor, MI**
M.S. in Computer Science and Engineering — *May 2015*
GPA: 3.96/4.0

Shanghai Jiao Tong University — **Shanghai, China**
B.S. in Electrical Engineering — *August 2013*
GPA: 3.85/4.0

University of Michigan — **Ann Arbor, MI**
B.S. in Computer Engineering — *May 2013*
GPA: 3.63/4.0

WORK EXPERIENCE

Alibaba Group — **Sunnyvale, CA**
Research Intern — *May - August 2019*

- Extended MLModelScope with automatic across-stack analysis capability.
- Used MLModelScope to benchmark and characterize public, MLPerf and AI Matrix models across systems of interest.
- Performed model/framework/system advising using the data collected, and explore its applicability in the Alibaba Cloud.

IBM Thomas J. Watson Research Center — **Yorktown Heights, NY**
Research Intern — *May - August 2018*

- Evaluated existing techniques for Deep Learning performance estimation on different models and systems, and understood the sources of inaccuracy.
- Developed an analysis tool that generates layer benchmarks, finds patterns within models, and performs performance prediction for Deep Learning models across hardware.

9th Programming and Tuning Massively Parallel Systems and AI School **Barcelona, Spain**
Teaching Assistant *July 2018*

- Designed GPU labs and projects for the summer school students.
- Advised the students during the summer school's hackathon.

IBM Thomas J. Watson Research Center **Yorktown Heights, NY**
Research Intern *May - August 2017*

- Developed **MLModelScope** a hardware/software agnostic platform for consistent benchmarking and analysis of Deep Learning inference at scale.
- Profiled and optimized the GPU-accelerated alternating least square(ALS) algorithm for Matrix Factorization.

University of Illinois Urbana-Champaign **Champaign, IL**
Lead Teaching Assistant for CS483 - Applied Parallel Programming *August - December 2016*

- Designed GPU labs, exams, and projects for a class of 200 students. Maintained the assignment and the project submission systems - **WebGPU** and **RAI**.

RECENT PROJECTS

XSP
- XSP is an across-stack profiling design that innovatively leverages distributed tracing to aggregate profile data from different profiling sources and construct a holistic and hierarchical view of DL model execution.
- XSP introduces a leveled and iterative measurement approach that accurately captures the latencies at all levels of the HW/SW stack despite the profiling overhead.
- We implement the design for GPUs and couple it with an automated analysis pipeline that enables systematic characterization and comparison.

Benanza
- We propose a "lower-bound" latency metric for DL models on GPUs based on the observation that the latency of a DL model is bounded by the latencies of the cuDNN and cuBLAS API calls corresponding to the model layers.
- Benanza is a sustainable and extensible benchmarking and analysis design that automatically generates micro-benchmarks given a set of models, computes their "lower-bound" latencies using the benchmark data, and informs optimizations of their execution on GPUs.

MLModelScope
- MLModelScope is a framework- and hardware-agnostic distributed platform for benchmarking and profiling DL models across datasets/frameworks/systems.
- MLModelScope proposes a specification to define DL model evaluations and techniques to provision the evaluation workflow using the user-specified HW/SW stack.
- MLModelScope is implemented as an open-source project with support for all major frameworks and hardware architectures.

DLBricks
- DLBricks is a composable benchmark generation design that reduces the effort of developing, maintaining, and running DL benchmarks on CPUs.
- DLBricks decomposes DL models into a set of unique runnable networks and constructs the original model's performance using the performance of the generated benchmarks.

TrIMS: Transparent and Isolated Model Sharing for DL Inference
- TrIMS is a generic memory sharing technique that enables constant data to be shared across processes or containers while still maintaining isolation between users.
- TrIMS mitigates the DL model loading overhead and increases the hardware resource utilization in inference by sharing models across all levels of the memory hierarchy in the cloud environment — GPU, CPU, local storage, and remote storage.

TOPS

- TOPS is a library of collectives expressed as matrix multiplication operations on Tensor Cores Units (TCU, specialized hardware for matrix multiplication).
- It is the first to broaden the class of algorithms expressible as TCU operations and show benefits of the mapping in terms of program simplicity, efficiency, and performance.
- We implemented reduction and scan using NVIDIA V100 Tensor Cores and achieved up to $100\times$ and $3\times$ speedup compared to state-of-the-art methods while decreasing the power consumption by up to 22% and 16% correspondingly.

PUBLICATIONS

1. **XSP: Across-Stack Profiling and Analysis of Machine Learning Models on GPUs (IPDPS'20, Best Paper Nomination)**
   Cheng Li*, *Abdul Dakkak**, *Jinjun Xiong*, *Wei Wei*, *Lingjie Xu*, *Wen-Mei Hwu*
2. **Benanza: Automatic uBenchmark Generation to Compute "Lower-bound" Latency and Inform Optimizations of Deep Learning Models on GPUs (IPDPS'20)**
   Cheng Li*, *Abdul Dakkak**, *Jinjun Xiong*, *Wen-Mei Hwu*
3. **DLBricks: Composable Benchmark Generation to Reduce Deep Learning Benchmarking Effort on CPUs (ICPE'20)**
   Cheng Li, *Abdul Dakkak*, *Jinjun Xiong*, *Wen-Mei Hwu*
4. **The Design and Implementation of a Scalable DL Benchmarking Platform (arXiv'19)**
   Cheng Li, *Abdul Dakkak*, *Jinjun Xiong*, *Wen-Mei Hwu*
5. **AI Matrix: A Deep Learning Benchmark for Alibaba Data Centers (arXiv'19)**
   Wei Zhang, *Wei Wei*, *Lingjie Xu*, *Lingling Jin*, ***Cheng Li***
6. **MLModelScope: Evaluate and Introspect Cognitive Pipelines (IEEE Services'19)**
   Cheng Li, *Abdul Dakkak*, *Jinjun Xiong*, *Wen-Mei Hwu*
7. **TrIMS: Transparent and Isolated Model Sharing for Low Latency Deep Learning Inference in Function as a Service Environments (IEEE CLOUD'19)**
   Abdul Dakkak, ***Cheng Li***, *Simon Garcia de Gonzalo*, *Jinjun Xiong*, *Wen-Mei Hwu*
8. **Accelerating Reduction and Scan Using Tensor Core Units (ICS'19)**
   Abdul Dakkak, ***Cheng Li***, *Jinjun Xiong*, *Isaac Gelado*, *Wen-Mei Hwu*
9. **Evaluating Characteristics of CUDA Communication Primitives on High-Bandwidth Interconnects (ICPE'19, Best Paper)**
   Carl Pearson, *Abdul Dakkak*, *Sarah Hashash*, ***Cheng Li***, *I-Hsin Chung*, *Jinjun Xiong*, *Wen-Mei Hwu*
10. **Accelerating Reduction Using Tensor Core Units (HPCaML'19)**
    Abdul Dakkak, ***Cheng Li***, *Jinjun Xiong*, *Wen-Mei Hwu*
11. **SCOPE: C3SR Systems Characterization and Benchmarking Framework (arXiv'18)**
    Carl Pearson, *Abdul Dakkak*, ***Cheng Li***, *Sarah Hashash*, *Jinjun Xiong*, *Wen-mei Hwu*
12. **Matrix Factorization on GPUs with Memory Optimization and Approximate Computing (ICPP'18)**
    Wei Tan, *Shiyu Chang*, *Liana Fong*, ***Cheng Li***, *Zijun Wang*, *LiangLiang Cao*
13. **RAI: A Scalable Project Submission System for Parallel Programming Courses (IPDPSW'17)**
    Abdul Dakkak, *Carl Pearson*, ***Cheng Li***, *Wen-mei Hwu*
14. **KLAP: Kernel Launch Aggregation and Promotion for Optimizing Dynamic Parallelism (MICRO'16)**
    Izzat El Hajj, *Juan Gomez-Luna*, ***Cheng Li***, *Li-Wen Chang*, *Dejan Milojicic*, *Wen-mei Hwu*
15. **DjiNN and Tonic: DNN as a Service and Its Implications for Future Warehouse Scale Comput-**

ers **(ISCA'15)**
*Johann Hauswald, Yiping Kang, Michael A. Laurenzano, Quan Chen, **Cheng Li**, Trevor Mudge, Ronald G. Dreslinski, Jason Mars, Lingjia Tang*

16. **Sirius: An Open End-to-End Voice and Vision Personal Assistant and Its Implications for Future Warehouse Scale Computers** **(ASPLOS'15)**
*Johann Hauswald, Michael A. Laurenzano, Yunqi Zhang, **Cheng Li**, Austin Rovinski, Arjun Khurana, Ronald G. Dreslinski, Trevor Mudge, Vinicius Petrucci1, Lingjia Tang, Jason Mars*

17. **Stochastic circuits for real-time image-processing applications** **(DAC'13)**
*Armin Alaghi, **Cheng Li**, John P. Hayes*

TALKS & POSTERS

Super Computing 2019 **Denver, CO**
Across-stack Profiling and Analysis of ML Models on GPUs *November 18, 2019*

Tutorial at IISWC 2019 **Orlando, FL**
Challenges and Solutions for End-to-End and Across Stack ML Benchmarking *November 3, 2019*

HotChips 2019 **Palo Alto, CA**
MLModelScope: Evaluate and Profile ML Models at Scale and Across Stack *August 18, 2019*

Tutorial at ISCA 2019 **Phoenix, AZ**
Benchmarking Deep Learning Systems *June 22, 2019*

Tutorial at ASPLOS 2019 **Providence, RI**
Benchmarking Deep Learning Systems *April 14, 2019*

NVIDIA GPU Technology Conference 2019 **San Jose, CA**
TOPS: Accelerating Reduction Using Tensor Core Units *March 22, 2019*

NVIDIA GPU Technology Conference 2019 **San Jose, CA**
TrIMS: Transparent and Isolated Model Sharing for Low Latency DL Inference *March 22, 2019*

NVIDIA GPU Technology Conference 2019 **San Jose, CA**
MLModelScope *March 22, 2019*

Super Computing 2018 **Dallas, TX**
MLModelScope *November 11, 2018*

IBM AI Research Week 2018 **Boston, MA**
MLModelScope *October 11, 2018*

NVIDIA GPU Technology Conference 2017 **San Jose, CA**
RAI: A Scalable Submission System for GPU Applications *March 22, 2017*

LANGUAGES

C/C++, Go, CUDA, Python, JavaScript, Bash, Mathematica

Chinese, English

MEMBERSHIP

IEEE, ACM, CRA-W (Computing Research Association-Women), WCS (Women in Computer Science)